



# RemindMe App

Team#1 - Final Presentation



# Agenda

1. Application Overview
2. Architecture
3. AP: Shared Sessions
4. AP: Logging
5. Experiments
6. Demo

# Application Overview






# Application Overview

Users can

- Sign up / sign in
- Create, read, update, and delete their reminders with information such as
  - Title
  - Description
  - Public or Not
- Public reminders can be viewed publicly
- Only the creator of the reminder can edit or delete the record; other users can view the reminder if it is set to public.


Architecture



# Software Stack

## Server 1

### Load Balancing



docker Swarm

NGINX


## Server 2

### backend

django

django  
REST  
framework

### Database

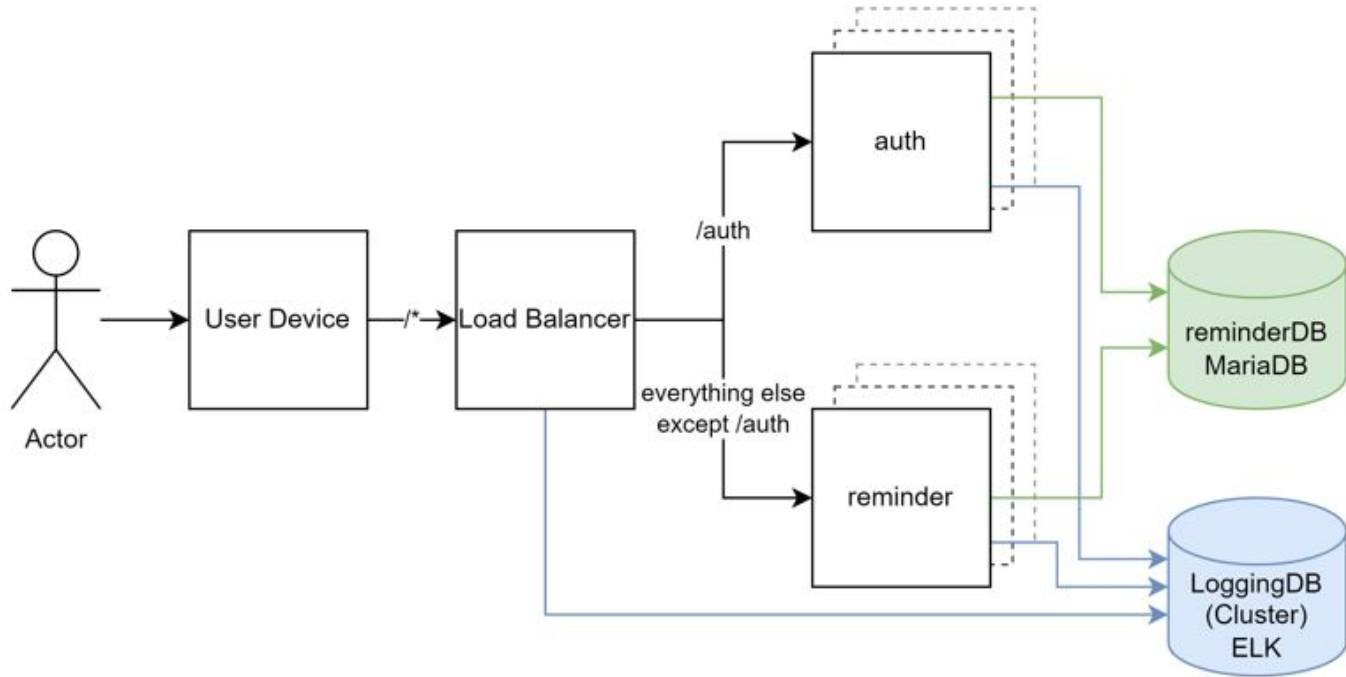


MariaDB



elasticsearch

# Overview





# Data Schema

## User:

- email
- username
- Password
  
- id
- is\_superuser
- first\_name
- last\_name
- ...

## Reminder

- id
- title
- content
- completed
- regisDate
- created\_by
- public



# AP: Shared Sessions





# AP: Shared Sessions

**Problem:** Many independent django instances must be aware of client's state (e.g. authentication state, user's application state (if applicable), ...)

**Solution:** Client-based session with JWT

- Tokens are signed by authentication service upon successful authentication
- May contain data such as username and expiration date
- Data can't be changed by the client without breaking the signature

⇒ Each reminder service-instance can verify tokens **independently** and **without accessing** the database (unlimited scaling)



# Typical Problems with JWT

**Statelessness:** Token can't be adapted to the application's state

⇒ Our application is designed to be stateless

**How to invalidate tokens:** Tokens are still usable even after logging out

- Remove token from client  
⇒ doesn't help if someone stole the token earlier
- Keep a database of blacklisted tokens  
⇒ defies the purpose of going client-based (but at least the database would be small)
- Keep expiration times short and get regular refresh tokens  
⇒ doesn't solve the problem, only shortens the time in which outdated tokens are valid  
⇒ additional load for authentication service



# Why not database-side or server-side session?

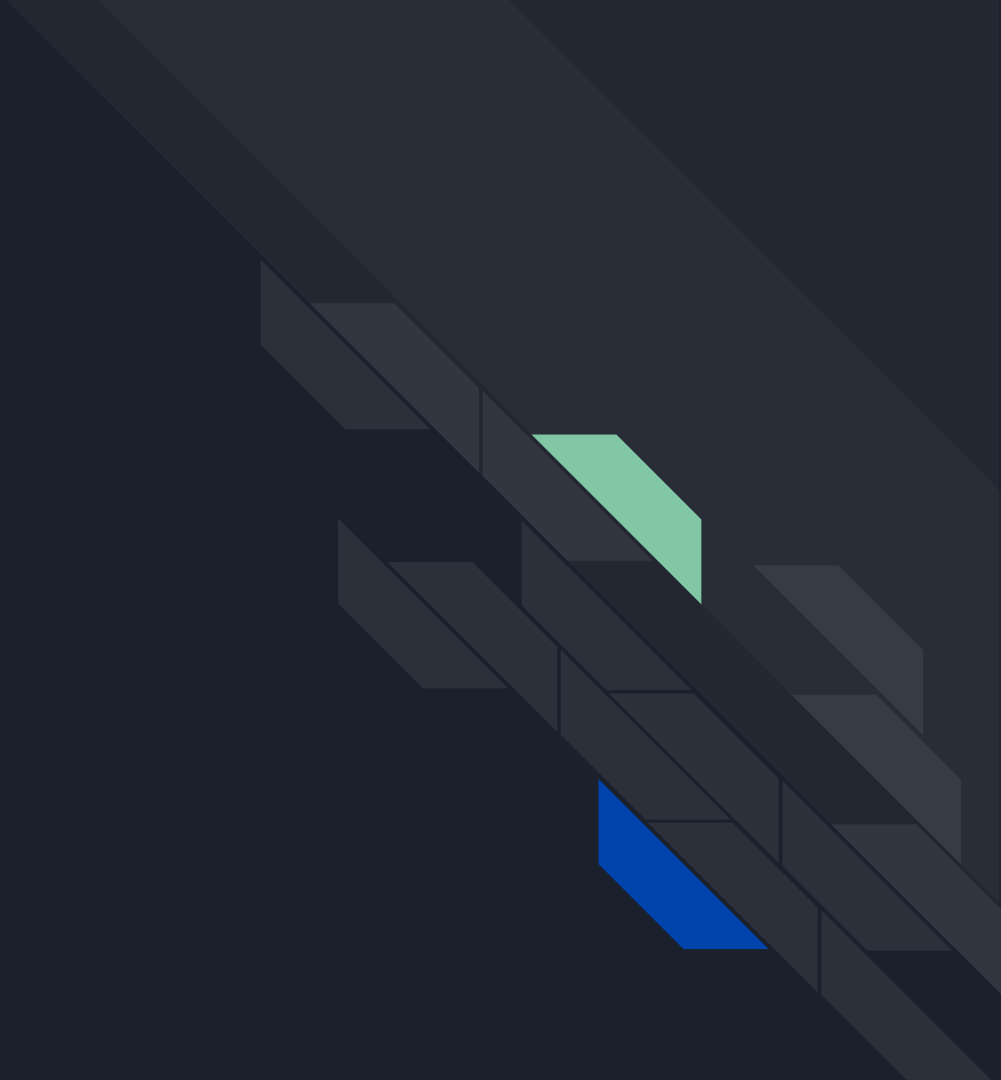
## Database-side session handling:

- Slower and poor scaling
- Doesn't work if the database breaks (single point of failure)

## Server-side session handling:

- Requests need to be forwarded to the same server instance or servers need to synchronize themselves
- Risk of data loss if an instance dies

AP: Logging





# AP: Logging - Overview

Why not simply use MariaDB again to store all logs messages?

- Setup would be the easiest way
- Would probably work

⇒ Log data is inherently different from business data

- Only create and read, never update; delete based on data retention policy
- Can make use of bulk operations
- Must be easily searchable and displayable for developers



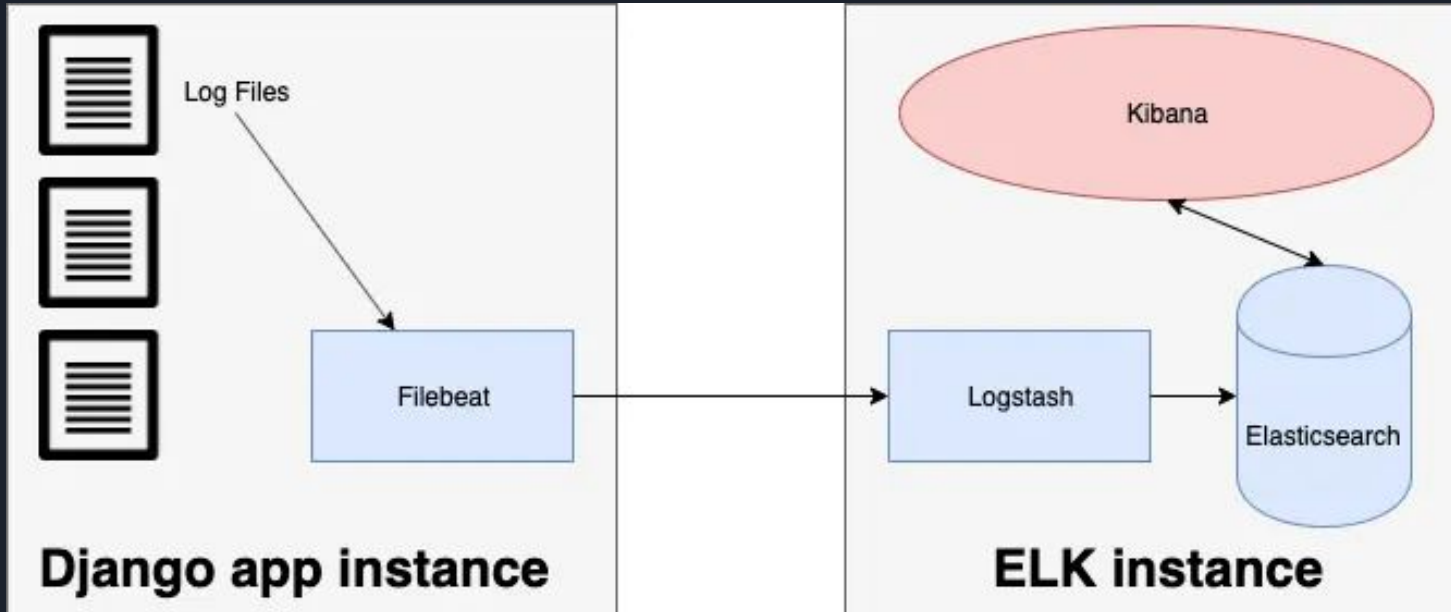
# AP: Logging - Proposed Solution

Proposed solution: Filebeat → Logstash → Elasticsearch → Kibana

1. Log messages are not sent to the database one-by-one, but collected locally first
2. **Filebeat** fetches local log files and ships them to Logstash
3. **Logstash** processes log messages and sends them to **Elasticsearch**
4. **Elasticsearch** stores log data in indices
5. **Kibana** displays log data from Elasticsearch indices in a dashboard

**Due to problems with the set-up, the solution is not fully implemented yet.**

# AP: Logging - Proposed Solution







# AP: Logging - Pros and Cons

## Pros

- Highly optimized and scalable for querying and indexing of log data
- Already includes the frontend dashboard
- Log files can be collected from any service and container

## Cons

- ELK stack requires a lot effort for maintenance and setup
- High resource requirements

# Experiments





# Experiments

## Short-term scope:

1. Load and performance testing with Apache JMeter
2. Resilience testing by turning of some of the docker nodes

## Long-term scope:

3. More physical machines → turn some machines off randomly
4. Load testing of ELK stack

Demo



Thank you

